

# CSE 242 Homework 4

Sampson Mao, samao@ucsc.edu

December 14, 2020

## 1 Kernel methods with noisy setting

### 1.1 (a)

```
[1]: import scipy.stats as stats
import numpy as np
np.random.seed(100)
plus_train = stats.multivariate_normal.rvs(mean=[1,1], cov=[[1,0],[0,1]],
↪size=1000)
minus_train = stats.multivariate_normal.rvs(mean=[-1,-1], cov=[[3,0],[0,3]],
↪size=1000)
train = np.concatenate((plus_train,minus_train))
plus_test = stats.multivariate_normal.rvs(mean=[1,1], cov=[[1,0],[0,1]],
↪size=1000)
minus_test = stats.multivariate_normal.rvs(mean=[-1,-1], cov=[[3,0],[0,3]],
↪size=1000)
test = np.concatenate((plus_test,minus_test))

train_labels = np.array([1] * 1000 + [-1] * 1000)
test_labels = np.array([1] * 1000 + [-1] * 1000)
```

```
[2]: def flip_labels(labels):
    to_minus_prob = np.concatenate((stats.uniform.rvs(size=1000), np.
↪ones(1000)))
    to_plus_prob = np.concatenate((np.ones(1000),stats.uniform.rvs(size=1000)))
    randomized_labels = labels.copy()
    randomized_labels[to_minus_prob <= 0.35] = -1
    randomized_labels[to_plus_prob <= 0.2] = 1
    return(randomized_labels)
```

```
[3]: from sklearn.svm import SVC

np.random.seed(2020)
a_accuracy = []
a_sums = []
for i in range(20):
    noisy_labels = flip_labels(train_labels)
```

```

a_sums.append(sum(noisy_labels))
svm_a = SVC(kernel="rbf", C=1, gamma=0.01)
svm_a.fit(train, noisy_labels)
a_accuracy.append(svm_a.score(test, test_labels))
np.mean(a_accuracy)

```

[3]: 0.8102499999999999

The average is around 0.81

## 1.2 (b)

```

[4]: from sklearn.neighbors import KNeighborsClassifier

np.random.seed(2020)
b_accuracy = []
b_sums = []
for i in range(20):
    noisy_labels = flip_labels(train_labels)
    KNN = KNeighborsClassifier(n_neighbors=19)
    KNN.fit(train, noisy_labels)
    fixed_labels = KNN.predict(train)
    b_sums.append(sum(fixed_labels))
    svm_b = SVC(kernel="rbf", C=1, gamma=0.01)
    svm_b.fit(train, fixed_labels)
    b_accuracy.append(svm_b.score(test, test_labels))
np.mean(b_accuracy)

```

[4]: 0.8528249999999999

```

[5]: np.transpose(b_accuracy)

```

```

[5]: array([0.8435, 0.8535, 0.8545, 0.8485, 0.845 , 0.851 , 0.861 , 0.8505,
          0.8495, 0.855 , 0.8395, 0.855 , 0.8525, 0.862 , 0.858 , 0.8575,
          0.858 , 0.8675, 0.8465, 0.848 ])

```

$k$  was chosen empirically and 19 seems to give the highest accuracy. It looks like there is a definite improvement, around 5%. Though we repeated the process 20 times to make sure the improvement is not by chance, it seems that each time is higher than 80% anyway.

## 1.3 (c)

```

[6]: from sklearn.cluster import KMeans

np.random.seed(2020)
c_accuracy = []
c_sums = []

```

```

for i in range(20):
    noisy_labels = flip_labels(train_labels)
    KM = KMeans(n_clusters=2)
    KM.fit(train)
    fixed_labels_KM = KM.predict(train)
    fixed_labels_KM[fixed_labels_KM == 0] = -1
    c_sums.append(sum(fixed_labels_KM))
    svm_c = SVC(kernel="rbf", C=1, gamma=0.01)
    svm_c.fit(train, fixed_labels_KM)
    c_accuracy.append(svm_c.score(test, test_labels))
np.mean(c_accuracy)

```

[6]: 0.39635

[7]: KM.cluster\_centers\_

[7]: array([[ 0.97198463, 0.91039327],  
 [-1.54659148, -1.46298677]])

$k$  was chosen to be 2 because we know that there are 2 clusters. There is a marked decrease in performance. One of the cluster centers (0.972, 0.910) was close to the mean we used to generate the data, but the other is not, which explains why. Also, increasing this parameter seems to make the accuracy worse. Overall, this method seems to make the data noisier and would not be recommended as a way to “fix” the labels.

## 2 Fairness

```

[8]: import sqlalchemy
engine = sqlalchemy.create_engine('sqlite:///compas.db')
inspector = sqlalchemy.inspect(engine)
inspector.get_table_names()

```

[8]: ['casearrest',  
 'charge',  
 'compas',  
 'jailhistory',  
 'people',  
 'prisonhistory',  
 'summary']

```

[9]: cursor = engine.execute('SELECT * FROM compas LIMIT 10')
cursor.fetchall()

```

[9]: [(1, 'miguel', 'hernandez', 56418, 58393, 64352, 'PRETRIAL', 'Risk and  
 Prescreen', 'Intake', 'Pretrial', 'Jail Inmate', 'Single', '2013-08-14  
 00:00:00.000000', 1, 'Low', 'Low', 7, 'Risk of Violence', -4.31, 1, 1),  
 (2, 'miguel', 'hernandez', 56418, 58393, 64352, 'PRETRIAL', 'Risk and

```

Prescreen', 'Intake', 'Pretrial', 'Jail Inmate', 'Single', '2013-08-14
00:00:00.000000', 1, 'Low', 'Low', 8, 'Risk of Recidivism', -2.78, 1, 1),
(3, 'miguel', 'hernandez', 56418, 58393, 64352, 'PRETRIAL', 'Risk and
Prescreen', 'Intake', 'Pretrial', 'Jail Inmate', 'Single', '2013-08-14
00:00:00.000000', 1, 'Low', 'Low', 18, 'Risk of Failure to Appear', 13, 1, 1),
(4, 'michael', 'ryan', 68603, 72042, 79669, 'PRETRIAL', 'Risk and Prescreen',
'Intake', 'Pretrial', 'Jail Inmate', 'Single', '2014-12-31 00:00:00.000000', 1,
'Low', 'Low', 7, 'Risk of Violence', -2.75, 2, 2),
(5, 'michael', 'ryan', 68603, 72042, 79669, 'PRETRIAL', 'Risk and Prescreen',
'Intake', 'Pretrial', 'Jail Inmate', 'Single', '2014-12-31 00:00:00.000000', 1,
'Low', 'Medium', 8, 'Risk of Recidivism', -0.34, 5, 2),
(6, 'michael', 'ryan', 68603, 72042, 79669, 'PRETRIAL', 'Risk and Prescreen',
'Intake', 'Pretrial', 'Jail Inmate', 'Single', '2014-12-31 00:00:00.000000', 1,
'Low', 'Low', 18, 'Risk of Failure to Appear', 16, 2, 2),
(7, 'kevon', 'dixon', 51601, 52796, 58102, 'PRETRIAL', 'Risk and Prescreen',
'Intake', 'Pretrial', 'Jail Inmate', 'Single', '2013-01-27 00:00:00.000000', 1,
'Low', 'Low', 7, 'Risk of Violence', -3.07, 1, 3),
(8, 'kevon', 'dixon', 51601, 52796, 58102, 'PRETRIAL', 'Risk and Prescreen',
'Intake', 'Pretrial', 'Jail Inmate', 'Single', '2013-01-27 00:00:00.000000', 1,
'Low', 'Low', 8, 'Risk of Recidivism', -0.76, 3, 3),
(9, 'kevon', 'dixon', 51601, 52796, 58102, 'PRETRIAL', 'Risk and Prescreen',
'Intake', 'Pretrial', 'Jail Inmate', 'Single', '2013-01-27 00:00:00.000000', 1,
'Low', 'Medium', 18, 'Risk of Failure to Appear', 25, 6, 3),
(10, 'ed', 'philo', 38864, 55421, 61042, 'PRETRIAL', 'Risk and Prescreen',
'Intake', 'Pretrial', 'Jail Inmate', 'Single', '2013-04-14 00:00:00.000000', 1,
'Low', 'Low', 7, 'Risk of Violence', -2.26, 3, 4)]

```

```

[10]: # Import the table "people" from the database into pandas
import pandas
people = pandas.read_sql('people', engine)
people.head()

```

```

[10]:   id      name  first  last  sex      race \
0    1  miguel hernandez  miguel  hernandez  Male      Other
1    2    michael ryan  michael    ryan  Male      Caucasian
2    3    kevon dixon   kevon    dixon  Male  African-American
3    4      ed philo     ed    philo  Male  African-American
4    5   marcu brown   marcu    brown  Male  African-American

      dob  age  age_cat  juv_fel_count  ...  r_offense_date \
0  1947-04-18   69  Greater than 45      0  ...      NaT
1  1985-02-06   31      25 - 45      0  ...      NaT
2  1982-01-22   34      25 - 45      0  ...  2013-07-05
3  1991-05-14   24  Less than 25      0  ...  2013-06-16
4  1993-01-21   23  Less than 25      0  ...      NaT

      r_charge_desc      r_jail_in      r_jail_out \

```

0		None	NaT	NaT
1		None	NaT	NaT
2	Felony Battery (Dom Strang)		NaT	NaT
3	Driving Under The Influence	2013-06-16 09:05:47	2013-06-16 07:18:55	
4		None	NaT	NaT

	is_violent_recid	num_vr_cases	vr_case_number	vr_charge_degree	\
0	0	None	None	None	
1	0	None	None	None	
2	1	None	13009779CF10A	(F3)	
3	0	None	None	None	
4	0	None	None	None	

	vr_offense_date	vr_charge_desc
0	NaT	None
1	NaT	None
2	2013-07-05	Felony Battery (Dom Strang)
3	NaT	None
4	NaT	None

[5 rows x 41 columns]

```
[11]: people.columns
```

```
[11]: Index(['id', 'name', 'first', 'last', 'sex', 'race', 'dob', 'age', 'age_cat',
        'juv_fel_count', 'juv_misd_count', 'juv_other_count',
        'compas_screening_date', 'decile_score', 'score_text', 'violent_recid',
        'priors_count', 'days_b_screening_arrest', 'c_jail_in', 'c_jail_out',
        'c_case_number', 'c_days_from_compas', 'c_arrest_date',
        'c_offense_date', 'c_charge_degree', 'c_charge_desc', 'is_recid',
        'num_r_cases', 'r_case_number', 'r_charge_degree', 'r_days_from_arrest',
        'r_offense_date', 'r_charge_desc', 'r_jail_in', 'r_jail_out',
        'is_violent_recid', 'num_vr_cases', 'vr_case_number',
        'vr_charge_degree', 'vr_offense_date', 'vr_charge_desc'],
        dtype='object')
```

```
[12]: # Import the table "compas" from the database into pandas
      # COMPAS is the risk assessment tool used in the article and is the subject of
      ↪ study
      compas = pandas.read_sql('compas', engine)
      compas.head()
```

	id	first	last	compas_person_id	compas_case_id	\
0	1	miguel	hernandez	56418	58393	
1	2	miguel	hernandez	56418	58393	
2	3	miguel	hernandez	56418	58393	
3	4	michael	ryan	68603	72042	

```

4    5    michael          ryan          68603          72042

      compas_assessment_id agency_text          scale_set assessment_reason \
0              64352      PRETRIAL Risk and Prescreen          Intake
1              64352      PRETRIAL Risk and Prescreen          Intake
2              64352      PRETRIAL Risk and Prescreen          Intake
3              79669      PRETRIAL Risk and Prescreen          Intake
4              79669      PRETRIAL Risk and Prescreen          Intake

      legal_status ... marital_status screening_date rec_supervision_level \
0      Pretrial ...      Single      2013-08-14          1
1      Pretrial ...      Single      2013-08-14          1
2      Pretrial ...      Single      2013-08-14          1
3      Pretrial ...      Single      2014-12-31          1
4      Pretrial ...      Single      2014-12-31          1

      rec_supervision_level_text score_text scale_id          type_of_assessment \
0              Low      Low      7          Risk of Violence
1              Low      Low      8          Risk of Recidivism
2              Low      Low      18 Risk of Failure to Appear
3              Low      Low      7          Risk of Violence
4              Low      Medium      8          Risk of Recidivism

      raw_score decile_score person_id
0          -4          1          1
1          -2          1          1
2          13          1          1
3          -2          2          2
4           0          5          2

```

[5 rows x 21 columns]

```
[13]: compas.columns
```

```
[13]: Index(['id', 'first', 'last', 'compas_person_id', 'compas_case_id',
        'compas_assessment_id', 'agency_text', 'scale_set', 'assessment_reason',
        'legal_status', 'custody_status', 'marital_status', 'screening_date',
        'rec_supervision_level', 'rec_supervision_level_text', 'score_text',
        'scale_id', 'type_of_assessment', 'raw_score', 'decile_score',
        'person_id'],
        dtype='object')
```

```
[14]: # Inner join the tables. Compare empirical data (times reoffended) with
      ↪ predicted (compas risk)

df = pandas.read_sql(''SELECT race,is_violent_recid,agency_text,compas.
      ↪score_text
```

```
FROM people JOIN compas ON
person_id = people.id'', engine)
df.head()
```

```
[14]:      race  is_violent_recid agency_text score_text
0      Other                0    PRETRIAL        Low
1      Other                0    PRETRIAL        Low
2      Other                0    PRETRIAL        Low
3  Caucasian                0    PRETRIAL        Low
4  Caucasian                0    PRETRIAL    Medium
```


```
[15]: relative_counts = pandas.DataFrame(
      {i: d.score_text.value_counts() / d.score_text.count()
       for i, d in df.groupby(['race', 'is_violent_recid'])})
relative_counts
```

```
[15]:      African-American      Asian      Caucasian \
      0      1      0      1      0      1
High      0.168173  0.290386  0.041667  0.166667  0.076904  0.142686
Low       0.576348  0.399707  0.839286  0.416667  0.742984  0.594724
Medium    0.255006  0.309907  0.119048  0.416667  0.179191  0.262590
N/A       0.000473      NaN      NaN      NaN      0.000922      NaN

      Hispanic      Native American      Other
      0      1      0      1      0      1
High  0.065251  0.159204      0.117117  0.277778  0.025402  0.191489
Low   0.780240  0.621891      0.711712  0.388889  0.855365  0.666667
Medium 0.154509  0.218905      0.171171  0.333333  0.118196  0.141844
N/A      NaN      NaN      NaN      NaN      0.001037      NaN
```

```
[16]: %matplotlib inline
relative_counts.plot.bar().legend(bbox_to_anchor = (1,1))
```

```
[16]: <matplotlib.legend.Legend at 0x7ffe74567350>
```



```
[17]: # this will be your prediction
df['scored_high'] = 0 + (df.score_text == 'High')
```

```
[18]: from statsmodels.formula.api import logit
#this will be the model you used to generate scores
# C() for categorical
model = logit('scored_high ~ C(race) + is_violent_recid + agency_text', df).
    ↪fit()
model.summary()
```

Optimization terminated successfully.  
 Current function value: 0.365167  
 Iterations 7

```
[18]: <class 'statsmodels.iolib.summary.Summary'>
      """
```

```

                                Logit Regression Results
=====
Dep. Variable:          scored_high    No. Observations:          37578
Model:                  Logit         Df Residuals:              37568
Method:                 MLE           Df Model:                  9
Date:                  Mon, 14 Dec 2020    Pseudo R-squ.:            0.04752
Time:                  13:19:34          Log-Likelihood:           -13722.
converged:              True            LL-Null:                  -14407.
Covariance Type:        nonrobust        LLR p-value:              3.681e-289
=====
=====
```



	coef	std err	z	P> z
[0.025      0.975]				
-----				
Intercept	-0.9779	0.230	-4.248	0.000
-1.429      -0.527				
C(race)[T.Asian]	-1.4142	0.343	-4.122	0.000
-2.087      -0.742				
C(race)[T.Caucasian]	-0.8806	0.038	-23.374	0.000
-0.954      -0.807				
C(race)[T.Hispanic]	-1.0405	0.069	-15.008	0.000
-1.176      -0.905				
C(race)[T.Native American]	-0.3605	0.256	-1.406	0.160
-0.863      0.142				
C(race)[T.Other]	-1.7172	0.119	-14.478	0.000
-1.950      -1.485				
agency_text[T.DRRD]	0.2336	0.294	0.795	0.427
-0.342      0.809				
agency_text[T.PRETRIAL]	-0.6537	0.231	-2.836	0.005
-1.105      -0.202				
agency_text[T.Probation]	-0.2184	0.241	-0.905	0.366
-0.692      0.255				
is_violent_recid	0.7479	0.045	16.515	0.000
0.659      0.837				
=====				
=====				
""""				

```
[19]: score = model.predict(df)
```

```
[20]: bw_score = score[(df["race"] == "African-American") | (df["race"] ==
    ↪ "Caucasian") ]
```

```
[21]: max_score = max(bw_score)
max_score
```

```
[21]: 0.500899933158618
```

```
[22]: min_score = min(bw_score)
min_score
```

```
[22]: 0.07501360138832201
```

## 2.1 (a) Independence

```
[23]: df_BW = df[(df["race"] == "Caucasian") | (df["race"] == "African-American")]
df_BW = df_BW.assign(R=(df_BW["scored_high"] >= np.
↳mean(df_BW["scored_high"]))*1)
```

```
[24]: P_W = sum((df_BW["race"] == "African-American") & (df_BW["R"] == 1))/
↳sum((df_BW["race"] == "African-American"))
P_B = sum((df_BW["race"] == "Caucasian") & (df_BW["R"] == 1))/
↳sum((df_BW["race"] == "Caucasian"))
P_W/P_B
```

```
[24]: 2.2336060238455064
```

```
[25]: np.mean(df_BW["scored_high"])
```

```
[25]: 0.1410753094585656
```

Using the empirical threshold of 0.141, the ratio evaluates to  $2.23 > 0.8$ . Therefore, the model violates disparate impact law

## 2.2 (b) Separation

```
[26]: def calc_TFPR(true, pred, n_thresh) :
x = np.linspace(min_score,max_score,n_thresh+1)
TPRs = []
FPRs = []
for thresh in x[0:n_thresh]:
converted_pred = (pred >= thresh)*1
TP = sum((true == 1) & (converted_pred == 1))
act_P = sum(true == 1)
TPR = TP/act_P
TPRs.append(TPR)
FP = sum((true == 0) & (converted_pred == 1))
act_N = sum(true == 0)
FPR = FP/act_N
FPRs.append(FPR)
df = pandas.DataFrame(zip(x[0:n_thresh], TPRs, FPRs),
columns=['Lower Threshold', 'TPR', 'FPR'])
return(df)
```

```
[27]: W_filter = df["race"] == "Caucasian"
W = df[W_filter]
W_score = score[W_filter]
TFPR_W = calc_TFPR(W["is_violent_recid"], W_score, 100)
TFPR_W
```

```
[27]:
```


	Lower Threshold	TPR	FPR
0	0.075014	1.0	1.000000
1	0.079272	1.0	0.036441
2	0.083531	1.0	0.036441
3	0.087790	1.0	0.036441
4	0.092049	1.0	0.036441
..	...	...	...
95	0.479606	0.0	0.000000
96	0.483864	0.0	0.000000
97	0.488123	0.0	0.000000
98	0.492382	0.0	0.000000
99	0.496641	0.0	0.000000

[100 rows x 3 columns]

```
[28]: import matplotlib.pyplot as plt

plt.plot(TFPR_W["FPR"], TFPR_W["TPR"])
plt.plot([0,1],[0,1],linestyle='-.')
plt.xlabel("FPR")
plt.ylabel("TPR")
plt.title("ROC Curve - Caucasian")
```

```
[28]: Text(0.5, 1.0, 'ROC Curve - Caucasian')
```



```
[29]: B_filter = df["race"] == "African-American"
      B = df[B_filter]
      B_score = score[B_filter]
      TFPR_B = calc_TFPR(B["is_violent_recid"], B_score, 100)
      TFPR_B
```

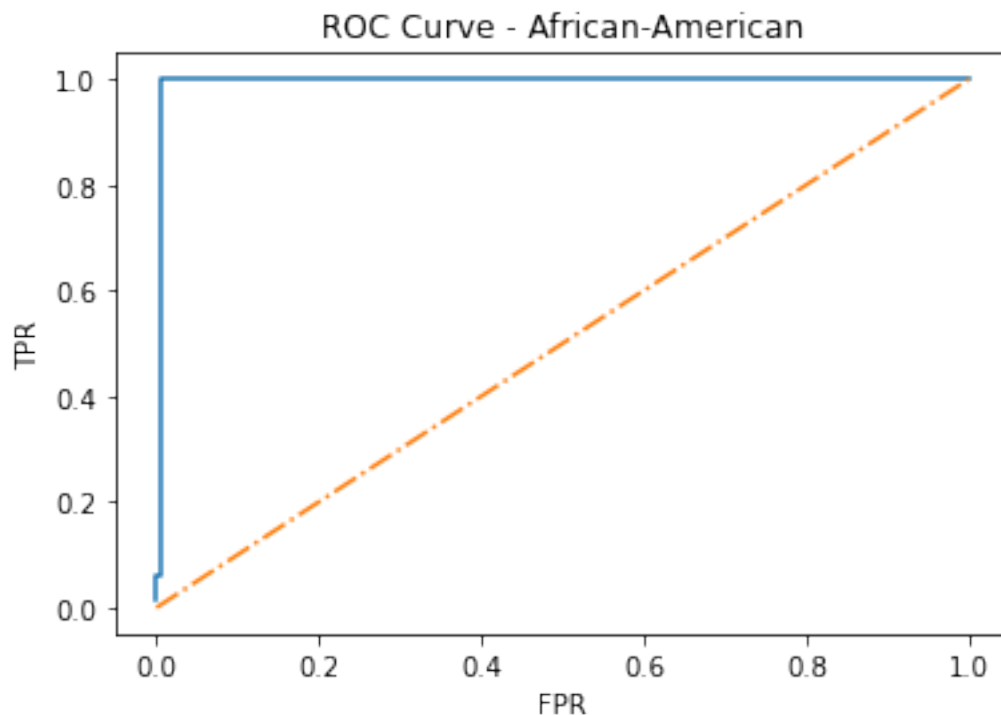
```
[29]:
```

	Lower Threshold	TPR	FPR
0	0.075014	1.000000	1.0
1	0.079272	1.000000	1.0
2	0.083531	1.000000	1.0
3	0.087790	1.000000	1.0
4	0.092049	1.000000	1.0
..	...	...	...
95	0.479606	0.013177	0.0
96	0.483864	0.013177	0.0
97	0.488123	0.013177	0.0
98	0.492382	0.013177	0.0
99	0.496641	0.013177	0.0

[100 rows x 3 columns]

```
[30]: plt.plot(TFPR_B["FPR"], TFPR_B["TPR"])
      plt.plot([0,1],[0,1],linestyle='-.')
      plt.xlabel("FPR")
      plt.ylabel("TPR")
      plt.title("ROC Curve - African-American")
```

```
[30]: Text(0.5, 1.0, 'ROC Curve - African-American')
```



### 2.3 (c) Achieve Separation by Post-Processing

The ideal TPR is 1 and the ideal FPR is 0. This would mean we should look for separation errors of 0. However, if the TPR for both groups is 0 and the FPR is 1, this also gives 0, and likewise if both the TPR and the FPR are 0. In terms of the ROC curve, we want the TPR and FPR corresponding to the upper left hand corner, with  $\text{TPR} = 1$ ,  $\text{FPR} \approx 0$ .

```
[31]: # oz stands for one zero
oz_TFPR_W = TFPR_W[(TFPR_W["TPR"] > 0) & (TFPR_W["FPR"] < 1)]
oz_TFPR_B = TFPR_B[(TFPR_B["TPR"] > 0) & (TFPR_B["FPR"] < 1)]
```

```
[32]: oz_TFPR_W.head()
```

```
[32]:
```

	Lower Threshold	TPR	FPR
1	0.079272	1.0	0.036441
2	0.083531	1.0	0.036441
3	0.087790	1.0	0.036441
4	0.092049	1.0	0.036441
5	0.096308	1.0	0.036441

```
[33]: oz_TFPR_B.head()
```

```
[33]:      Lower Threshold  TPR      FPR
      21      0.164450  1.0  0.041645
      22      0.168709  1.0  0.041645
      23      0.172967  1.0  0.041645
      24      0.177226  1.0  0.041645
      25      0.181485  1.0  0.041645
```

```
[34]: TPR_term = np.abs(np.subtract.outer(list(oz_TFPR_W["TPR"]),
      ↪list(oz_TFPR_B["TPR"])))
      FPR_term = np.abs(np.subtract.outer(list(oz_TFPR_W["FPR"]),
      ↪list(oz_TFPR_B["FPR"])))
      sep_error = TPR_term + FPR_term
```

```
[35]: x = oz_TFPR_W["Lower Threshold"]
      y = oz_TFPR_B["Lower Threshold"]
      coords = np.argwhere(sep_error == np.min(sep_error))
      pairs = [[x[w],y[b]] for w, b in zip(coords[:,0], coords[:,1])]
```

```
[36]: pairs = pandas.DataFrame(pairs, columns=["Caucasian", "African-American"])
      pairs
```

```
[36]:      Caucasian  African-American
      0      0.109085      0.185744
      1      0.109085      0.190003
      2      0.109085      0.194262
      3      0.109085      0.198521
      4      0.109085      0.202780
      5      0.113343      0.185744
      6      0.113343      0.190003
      7      0.113343      0.194262
      8      0.113343      0.198521
      9      0.113343      0.202780
     10      0.117602      0.185744
     11      0.117602      0.190003
     12      0.117602      0.194262
     13      0.117602      0.198521
     14      0.117602      0.202780
     15      0.121861      0.185744
     16      0.121861      0.190003
     17      0.121861      0.194262
     18      0.121861      0.198521
     19      0.121861      0.202780
     20      0.126120      0.185744
     21      0.126120      0.190003
     22      0.126120      0.194262
     23      0.126120      0.198521
     24      0.126120      0.202780
```

25	0.130379	0.185744
26	0.130379	0.190003
27	0.130379	0.194262
28	0.130379	0.198521
29	0.130379	0.202780

The table above are pairs of thresholds that would give the lowest separation error while maximizing TPR and minimizing FPR. The first column is the threshold for Caucasians and the second is for African-American. We can see from the table that the threshold is lower for Caucasians than African-Americans. We can also see this in the mean and median below. Recall that the score, the output from logistic regression, is the predicted probability of scoring high. Our goal is to look for bias in the COMPAS scoring system. We can see that this model implies that African-Americans more likely to get a high COMPAS score. Therefore, it suggests that there is bias in the COMPAS score, as race makes a big difference in the model.

```
[37]: pairs.describe()
```

```
[37]:
```

	Caucasian	African-American
count	30.000000	30.000000
mean	0.119732	0.194262
std	0.007398	0.006126
min	0.109085	0.185744
25%	0.113343	0.190003
50%	0.119732	0.194262
75%	0.126120	0.198521
max	0.130379	0.202780